

Data Frames

R Data Objects

Gaston Sanchez

CC BY-NC-SA 4.0

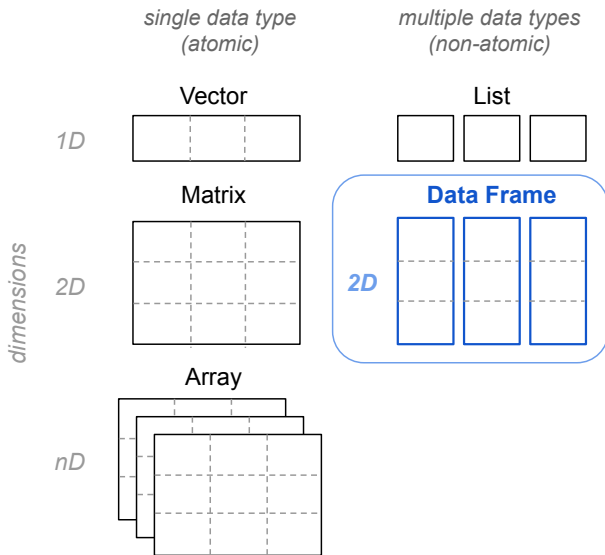
STAT 33B, Fall 2025

About

In this slides we introduce R data objects of class `data.frame`, which provide a nice tabular structure to work with.

Also, we discuss several functions for how to work with data frames from a classic perspective (as opposed to the alternative paradigm from a collection of non-native packages known as "tidyverse").

Basic Data Objects in R



About Data Frames

- ▶ R `data.frame`'s are internally stored as lists.
- ▶ At the same time, an R `data.frame` behaves like a 2-dimensional object (like a matrix).
- ▶ This means that you can manipulate a `data.frame` either like a list but also like a matrix.
- ▶ From the list point of view, each element of a `data.frame` is a column.
- ▶ A column typically corresponds to an atomic structure (vector or factor), but they can also correspond to a list (this is rare but it's possible).

Inspection Functions

Function	Description
<code>str()</code>	Structure
<code>head()</code>	First <code>n</code> rows
<code>tail()</code>	Last <code>n</code> rows
<code>dim()</code>	Dimensions (<code>#</code> rows, <code>#</code> cols)
<code>nrow()</code>	Number of rows
<code>ncol()</code>	Number of columns
<code>names()</code>	Vector of column names
<code>colnames()</code>	Vector of column names
<code>rownames()</code>	Vector of row names
<code>dimnames()</code>	List of row and column names
<code>summary()</code>	Descriptive statistics

Inspection Fuctions

A few rows of `airquality`

	Ozone	Solar.R	Wind	Temp	Month	Day
1	41	190	7.4	67	5	1
2	36	118	8.0	72	5	2
3	12	149	12.6	74	5	3
4	18	313	11.5	62	5	4
5	NA	NA	14.3	56	5	5
6	28	NA	14.9	66	5	6

Inspection Functions

```
# display structure
```

```
str(airquality)
```

```
# display structure (showing few elements)
```

```
str(airquality, vec.len = 1)
```

```
# first n rows
```

```
head(airquality, n = 5)
```

```
# last n rows
```

```
tail(airquality, n = 5)
```

```
# column summaries
```

```
summary(airquality)
```

Inspection Functions

```
# memory size  
object.size(airquality)
```

```
# attributes  
attributes(airquality)
```

```
# data frame dimensions  
dim(airquality)
```

```
# number of rows  
nrow(airquality)
```

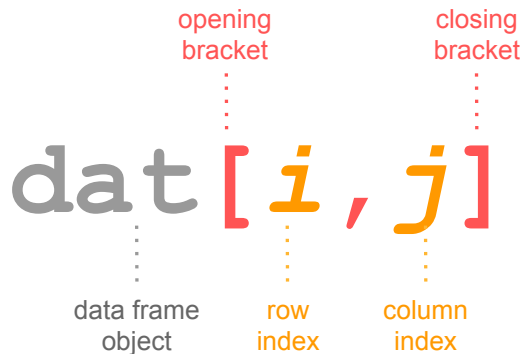
```
# number of columns  
ncol(airquality)
```


Inspection Functions

```
# row names  
rownames(airquality)  
  
# column names  
colnames(airquality)  
  
# column names  
names(airquality)  
  
# object class  
class(airquality)  
  
# check if object is data.frame  
is.data.frame(airquality)  
  
# data.frame is also a list  
is.list(airquality)
```


Subsetting and Indexing

Bracket Notation




Cells

`dat[2,2]`




one single
cell

`dat[2:3,2:3]`



consecutive
cells


`dat[c(1,3,5),c(2,4)]`




separate
cells

Cells


`dat[-2,-2]`



`dat[-(2:3),
-(2:3)]`



`dat[-c(1,3),
-c(2,4)]`



Columns

`dat[,3]`

one single
column

`dat[,1:3]`

consecutive
columns

`dat[,c(1,3)]`

separate
columns

Columns

`dat[, -3]`

one single
column

`dat[, -(1:3)]`

consecutive
columns

`dat[, -c(1,3)]`

separate
columns

Columns

```
# column Ozone  
airquality[ , "Ozone"]  
  
# columns Wind and Temp  
airquality[ , c("Wind", "Temp")]
```


Columns

You can use argument `drop = FALSE` when selecting one column, in order to keep the second dimension of the returned output

```
# first column (as a one-column data frame)
```

```
airquality[ , 1, drop=FALSE]
```

```
# column Ozone (as a one-column data frame)
```

```
airquality[ , "Ozone", drop=FALSE]
```

Rows

`dat[1,]`

one single
row

`dat[2:4,]`

consecutive
rows

`dat[c(1,4),]`

separate
rows

Rows

`dat[-1,]`

one single
row

`dat[-(2:4),]`

consecutive
rows

`dat[-c(1,4),]`

separate
rows

Bracket Notation

single brackets

`dat [: , name(s)]`

empty row index *column name(s)*

`dat $name`

double brackets

`dat [[j]]`

single column index
(integer, name)

Dollar Notation

```
# column Ozone  
airquality$Ozone
```

```
# equivalently  
airquality$"Ozone"
```

```
# equivalently  
airquality$'Ozone'
```

Equivalent Calls: brackets and \$

```
airquality[ , "Wind"]
```

```
airquality[ , "Wind", drop=FALSE]
```

```
airquality["Wind"]
```

```
airquality[["Wind"]]
```

```
airquality$Wind
```

Although the above commands are “equivalent” in the sense that they all allow you to select the data in column Wind, keep in mind that there are differences in terms of output format, and performance.

Modifying Data Frames

Modifying Data Frames

We've seen how to access the elements from a `"data.frame"`. Pretty much the same operations to extract the contents can be used to modify them.

We'll use a toy data set to see how we can modify the content of a `data.frame`.

Toy Data Frame


	name	gender	height	weight
1	Anakin	male	1.88	84
2	Padme	female	1.65	45
3	Luke	male	1.72	77
4	Leia	female	1.50	49

Toy Data Frame

```
# creating a data frame  
dat <- data.frame(  
  name = c('Anakin', 'Padme', 'Luke', 'Leia'),  
  gender = c('male', 'female', 'male', 'female'),  
  height = c(1.88, 1.65, 1.72, 1.50),  
  weight = c(84, 45, 77, 49)  
)
```

Modifying elements via index values

Modifying Elements via Index Values



Modifying Rows

```
dat[i, ] <- x
```


one single
row

consecutive
rows

separated
rows

Modifying Columns

`dat[,j] <- x`


one single
column

consecutive
columns


separated
columns

Modifying Cells


`dat` [*i*, *j*] \leftarrow *x*



one single
cell



consecutive
cells



separated
cells

Example: modifying dat

```
# affecting cell 1,1  
dat[1,1] <- 'ANAKIN'  
  
# gender as factor  
dat[,2] = factor(c('male',  
  'female', 'male', 'female'))
```


More options to modify columns

Example: modifying columns of dat


```
# name in upper case  
dat[, "name"] <- toupper(dat[, "name"])  
  
# height in cms  
dat$height <- dat$height * 100  
  
# weight in pounds  
dat[['weight']] <- dat[['weight']] * 2.20462
```

Example: logical subsetting

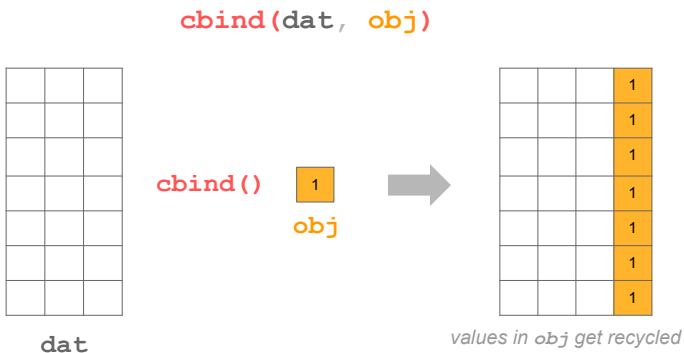
```
# change male names  
dat[dat$gender == 'male', 'name'] <- c('Anakin', 'Luke')  
  
# multiple conditions  
conds <- (dat$gender == 'female' & dat$height > 160)  
  
dat[conds, 'name'] <- "PADME"
```

Adding new elements

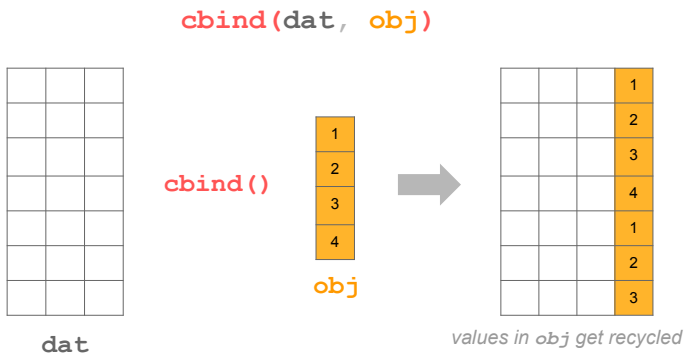
Adding columns with cbind()




Adding columns with cbind()



Adding columns with cbind()



Adding columns with cbind()



Adding columns with cbind()

```
# adding constant
```

```
dat = cbind(dat, constant = "force")
```

```
# add random column
```

```
dat = cbind(dat, random = runif(nrow(dat)))
```

Other ways to add columns

Equivalent ways to add a new column to data frame `dat`:

```
# add random column  
dat$new <- runif(nrow(dat))
```

```
# add random column  
dat[["new"]] <- runif(nrow(dat))
```

```
# add random column  
dat["new"] <- runif(nrow(dat))
```

Combining `dat[]` and `cbind()`

```
# number of columns
last = ncol(dat)
# increasing column indices
tmp = c(last+1, last+2)

# add new columns
dat[tmp] <- cbind(
  1:nrow(dat),
  rnorm(nrow(dat)))
```

Moving and Removing Columns

Moving columns

The common approach to move columns is to define a vector with the column names in the desired order, and then redefine the current data frame

```
# rearranged column names
rename = c('name', 'weight', 'height', "gender")

# moving columns
dat = dat[,rename]

# equivalently (list syntax)
dat[rename]
```

Removing a single column

To remove a column, you can indicate its name, and then *nullify it*

```
# equivalent calls
```

```
dat$weight <- NULL
```

```
dat[["weight"]] <- NULL
```

```
dat["weight"] <- NULL
```

Removing various columns

To remove various columns, you can indicate their names (via a character vector) or their column indices (via a numeric vector) and then *nullify them*

```
# equivalent calls
```

```
dat[,c("weight", "height")] <- NULL
```

```
dat[,c("weight", "height")] <- NULL
```

```
dat[,c(3, 4)] <- NULL
```

Sorting Elements

Sorting Elements

Some sorting functions:

- ▶ `sort()`
- ▶ `order()`

```
# Wind and Temp values  
# in Month 5, ordered by Wind  
wind5 = with(airquality,  
order(Wind[Month == 5]))  
  
airquality[wind5,  
c('Wind', 'Temp')]
```

Merging Tables


Merging Data Frames with `merge()`

Merging two or more data tables is another frequent type of operation. This can be done using the `merge()` function.

The behavior of `merge()` depends on a combination of several arguments. We'll see some of the frequent scenarios.

Merging: Simplest case

`merge(x, y, by = id)`



Same `id` columns (with same values) in both data frames.

In this case we could also use `cbind()`

Merging: Less simple case

`merge(x, y, by = id)`

id	x1	x2		id	y1	y2		id	x1	x2	y1	y2
a			&	d	4	8	→	a			1	5
b				c	3	7		b			2	6
c				b	2	6		c			3	7
d				a	1	5		d			4	8


Same `id` columns in both data frames.

Same `id` values but in different order.

Data frames merging (case 1)

Merging: Default merging

`merge(x, y, by = id, all = FALSE)`



same `id` columns in both data frames.

`id` in `y` is a subset of `id` in `x`

Merging `all = TRUE`

`merge(x, y, by = id, all = TRUE)`

id	x1	x2		id	y1	y2		id	x1	x2	y1	y2
a			&	a			→	a				
b				c				b			NA	NA
c								c				
d								d			NA	NA

same `id` columns in both data frames.

`id` in `y` is a subset of `id` in `x`

Merging `all.x = TRUE`

```
merge(x, y, by = id, all.x = TRUE)
```


id	x1	x2		id	y1	y2		id	x1	x2	y1	y2
a			&	a			➡	a				
b				c				b			NA	NA
c								c				
d								d			NA	NA

same `id` columns in both data frames.

`id` in `y` is a subset of `id` in `x`

Merging `all.y = TRUE`

```
merge(x, y, by = id, all.y = TRUE)
```




same `id` columns in both data frames.

`id` in `y` is a subset of `id` in `x`

Data frames merging (case 2)

Merging: Default all = FALSE


```
merge(x, y, by.x = idx, by.y = idy)
```



Different `id` columns in the data frames.

Merging: `all = TRUE`


```
merge(x, y, by.x = idx, by.y = idy,  
      all = TRUE)
```



Different `id` columns in the data frames.

Merging: `all.x = TRUE`


```
merge(x, y, by.x = idx, by.y = idy,  
      all.x = TRUE)
```



Different `id` columns in the data frames.

Merging: `all.y = TRUE`

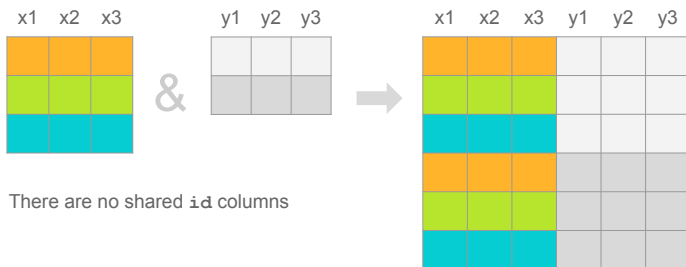
```
merge(x, y, by.x = idx, by.y = idy,  
      all.y = TRUE)
```



Different `id` columns in the data frames.

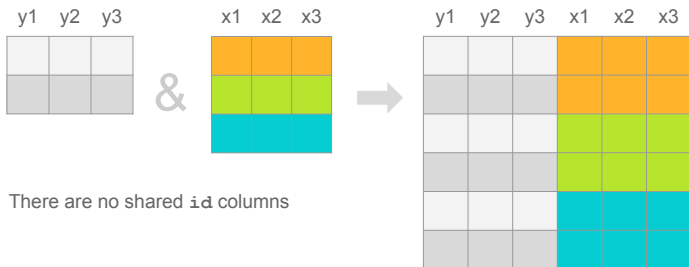
Merging with no matching column names

merge (**x**, **y**)



Merging with no matching column names

merge (*y*, *x*)



Tidyverse Approach

Later on we'll talk about another different approach for working with `data.frames` and similar tabular structures via the ecosystem of packages known as the **"tidyverse"**