

# Intro to Functions

## Programming Structures

Gaston Sanchez

CC BY-NC-SA 4.0

STAT 33B, Fall 2025

# About

In this part, we describe:

- ▶ the syntax for creating functions in R,
- ▶ the parts of an R function,
- ▶ various aspects about the arguments of functions

# Functions

# Writing Functions

- ▶ You can create functions with the function `function()`
- ▶ The arguments (inputs) go inside parenthesis, separated by commas.
- ▶ The code of the function is surrounded by braces (i.e. an *R expression*).

```
# example
square <- function(x) {
  x^2
}
```

# Writing Functions

Curly braces are optional if the body is a single expression.

```
square <- function(x) x^2
```

If the body of a function is a compound expression we have to use braces:

```
sum_sqr <- function(x, y) {  
  xy_sum <- x + y  
  xy_ssqr <- (xy_sum)^2  
  list(sum = xy_sum,  
        sumsqr = xy_ssqr)  
}
```

```
sum_sqr(3, 5)
```

```
$sum  
[1] 8
```

# Writing Functions

Once defined, functions can be used in other function definitions:

```
sum_of_squares <- function(x) {  
  sum(square(x))  
}
```

```
sum_of_squares(1:5)
```

```
[1] 55
```

# Nested Functions

We can also define a function inside another function:

```
sum_of_squares <- function(x) {  
  square <- function(x) {  
    x^2  
  }  
  sum(square(x))  
}
```

```
sum_of_squares(1:5)
```

```
[1] 55
```

# Anatomy of a function

Functions with a body consisting of a simple expression can be written with no braces (i.e. in one single line):

```
square <- function(x) x^2
```

```
square(5)
```

```
[1] 25
```



# Anatomy of a function

## Conceptual structure of a function

```
function_name <- function(arg1, arg2, etc) {  
  expression_1  
  expression_2  
  ...  
  expression_n  
}
```

# Anatomy of a function

- ▶ Generally we will assign the function to a name
- ▶ A function takes one or more inputs (or none), known as *arguments*
- ▶ The expressions forming the operations comprise the **body** of the function
- ▶ A function made of a simple expression doesn't require braces
- ▶ Functions return a single value

# Anatomy of a function

```
square <- function(x) {  
  x^2  
}
```

- ▶ the function's name is "square"
- ▶ it has one argument  $x$
- ▶ the function's body consists of one simple expression
- ▶ it returns the value  $x^2$

# Function Names

## Different ways to name functions

- ▶ `square()`
- ▶ `squ_are()`
- ▶ `s.square()`
- ▶ `Square()`
- ▶ `.square()`: a function that starts with a dot is a valid name, but the function will be a *hidden* function.

# Function names

## Invalid names

- ▶ `5square()`: cannot begin with a number
- ▶ `_square()`: cannot begin with an underscore
- ▶ `squ-are()`: cannot use hyphenated names

# Output of a Function

# Function Output

The value (i.e. output) of a function can be established in two ways:

- ▶ As the last evaluated simple expression (in the body)
- ▶ An explicitly **returned** value via `return()`

Recall that:

- ▶ The body of a function is an expression
- ▶ Remember that every expression has a value
- ▶ Hence every function has a value

# Function Output

Every function has a value (i.e. output)

```
cm2in <- function(x) {  
  x * 0.3937    # processing and output  
}
```

Recall that every expression has a value: the value of the last statement that is evaluated; in this case is `x * 0.3937`



# Function Output

Many useRs prefer to explicitly use a `return()` statement

```
cm2in <- function(x) {  
  y <- x * 0.3937    # processing  
  return(y)          # output  
}
```

# Function Output

Many amateur useRs like to use a `print()` statement (but you should avoid it)

```
cm2in <- function(x) {  
  y <- x * 0.3937    # processing  
  print(y)           # output  
}
```

Note: Using `print()` to specify the output of a function could work in most cases. But it largely depends on the object that is printed. Recall that `print()` is not a single function but a method—there are multiple flavors of `print()`. So to play safe, it's better to use `return()` than `print()`

## `return()` versus `print()`

- ▶ The function `print()` is a **generic** method in R.
- ▶ This means that `print()` has a different behavior depending on its input
- ▶ Unless you want to print intermediate results while the function is being executed, there is no need to return the output via `print()`

# Function Output

Keep in mind that depending on what's returned or what's the last evaluated expression, just calling a function might not print anything:

```
cm2in <- function(x) {  
  y <- x * 0.3937    # processing  
}
```

```
cm2in(5)
```

Note: the code of the function works, and the function has a value. But in this case the value is an assignment command, NOT an (implicit) printing command.

# Function Output

Here we call the function and assign it to an object.

```
cm2in <- function(x) {  
  y <- x * 0.3937    # processing  
}  
  
z = cm2in(5)  
z
```

```
[1] 1.9685
```

The last evaluated expression of the function has the same value as in the preceding slide. However, the very last command is a printing command which allows you to see the output.

# The return() command

`return()` can be useful when the output may be obtained in the middle of the function's body

```
plus_minus <- function(x, y, add = TRUE) {  
  if (add) {  
    return(x + y)  
  } else {  
    return(x - y)  
  }  
}
```

```
plus_minus(2, 3, add = TRUE)  
plus_minus(2, 3, add = FALSE)
```

## Return statement

Likewise, to exit the function and return a result early, use `return()`:

```
square <- function(x) {  
  if (!is.numeric(x)) {  
    return(NA)  
  }  
  x^2  
}  
square(6)
```

```
[1] 36
```

```
square("hi")
```

```
[1] NA
```

# Function Arguments



# Function Arguments

Functions can have any number of arguments (even zero arguments)

```
# function with 2 arguments
```

```
add <- function(x, y) x + y
```

```
# function with no arguments
```

```
hi <- function() print("Hi there!")
```

```
hi()
```

```
[1] "Hi there!"
```

# Arguments

Arguments can have default values (highly recommended!)

```
hey <- function(x = "") {  
  cat("Hey", x, "\nHow is it going?")  
}
```

```
hey()
```

Hey

How is it going?

```
hey("Gaston")
```

Hey Gaston

How is it going?

## Arguments with no default values

If you specify an argument with no default value, you must give it a value everytime you call the function, otherwise you'll get an error:

```
cm2in <- function(x) {  
  x * 0.3937  
}  
  
cm2in()
```

Error in cm2in(): argument "x" is missing, with no default

## Arguments with no default values

Sometimes we don't want to give default values, but we also don't want to cause an error. We can use `missing()` to see if an argument is missing:

```
abc <- function(a, b, c = 3) {  
  if (missing(b)) {  
    return((a * 2) + c)  
  } else {  
    return((a * b) + c)  
  }  
}
```

```
abc(1)
```

```
[1] 5
```

```
abc(1, 4)
```

```
[1] 7
```

# Arguments with no default values

You can also set an argument value to **NULL** if you don't want to specify a default value:

```
abc <- function(a, b, c = 3, d = NULL) {  
  if (is.null(d)) {  
    return((a * b) + c)  
  } else {  
    return((a * b) + (c * d))  
  }  
}
```

```
abc(1, 2)
```

```
[1] 5
```

```
abc(1, 2, 3, 4)
```

```
[1] 14
```

# More about function arguments

Arguments of functions can be:

▶ positional

```
# x and y are positional arguments  
plus <- function(x, y) x + y
```

▶ named

```
# x and y are named arguments  
plus <- function(x = 1, y = 1) x + y
```

# Argument Matching

```
# normal distribution
normal_distrib <- function(x, mu = 0, sigma = 1) {
  constant <- 1 / (sigma * sqrt(2*pi))
  constant * exp(-((x - mu)^2) / (2 * sigma^2))
}

normal_distrib(2)
normal_distrib(2, sigma = 3, mu = 1)
normal_distrib(mu = 1, sigma = 3, 2)
normal_distrib(mu = 1, 2, sigma = 3)
```

# Argument Matching

R is “smart” enough in doing pattern matching with arguments’ names (not recommended though)

```
normal_distrib(2)
```

```
[1] 0.05399097
```

```
normal_distrib(2, m = 0, s = 1)
```

```
[1] 0.05399097
```

```
normal_distrib(2, sig = 1, m = 0)
```

```
[1] 0.05399097
```



# Lazy Evaluation

In R, function arguments are **lazily evaluated**: they're only evaluated if needed.

For example, this code doesn't cause any problems because `x` is never used:

```
toss <- function(x) {  
  sample(c("heads", "tails"), size = 1)  
}  
  
toss()
```

# The dots parameter

The **dots** parameter `...` accepts any number of arguments, and it is often used to forward arguments to another function.

For example:

```
# Mean function with tolerance:
mean_tol <- function(x, tol, ...) {
  mean(x[x > tol], ...)
}

mean_tol(c(1, 3, 5, 0.01, 0.2, NA), 0.5)
mean_tol(c(1, 3, 5, 0.01, 0.2, NA), 0.5, na.rm = TRUE)
mean(c(1, 3, 5, 10))
```

# The dots parameter

You can access elements of `...` with the `...elt()` function:

```
hey <- function(x, ...) {  
  ...elt(2)  
  x + ...elt(1)  
}
```

```
hey(3, 5, message("hi"))
```

```
hi
```

```
[1] 8
```

This only evaluates the argument you accessed.

# The dots parameter

You can convert ... to a list with the `list()` function:

```
hey <- function(...) list(...)
```

```
hey(hi = 1, 3, 4)
```

```
$hi
```

```
[1] 1
```

```
[[2]]
```

```
[1] 3
```

```
[[3]]
```

```
[1] 4
```

This evaluates all of the arguments.

# Conditions

# Conditions

There are three main functions for generating warnings and errors:

- ▶ **message()**: to print an informative message
- ▶ **warning()**: to raise a warning message (without stopping execution)
- ▶ **stop()**: to stop execution raising an error

# Stop Execution

Use `stop()` to stop execution of a function (raising an error)

```
meansd <- function(x, na.rm = FALSE) {  
  if (!is.numeric(x)) {  
    stop("input must be numeric")  
  }  
  # output  
  c(mean = mean(x, na.rm = na.rm),  
    sd = sd(x, na.rm = na.rm))  
}
```

# Stop Execution

```
# ok
```

```
meansd(c(4, 5, 3, 1, 2))
```

```
      mean      sd  
3.000000 1.581139
```

```
# this causes an error
```

```
meansd(c('a', 'b', 'c'))
```

```
Error in meansd(c("a", "b", "c")): input must be numeric
```



# Warning Messages

A `warning()` is useful when we don't want to stop the execution, but we still want to show potential problems

```
meansd <- function(x, na.rm = FALSE) {  
  if (!is.numeric(x)) {  
    warning("non-numeric input coerced to numeric")  
    x <- as.numeric(x)  
  }  
  # output  
  c(mean = mean(x, na.rm = na.rm),  
    sd = sd(x, na.rm = na.rm))  
}
```

## Warning Message

```
# ok  
meansd(c(4, 5, 3, 1, 2))
```

```
      mean      sd  
3.000000 1.581139
```

```
# this causes a warning  
meansd(c(TRUE, FALSE, TRUE, FALSE))
```

```
Warning in meansd(c(TRUE, FALSE, TRUE, FALSE)):  
non-numeric input coerced to numeric
```

```
      mean      sd  
0.5000000 0.5773503
```

# Generic Messages

Use `message()` to display a generic message that is not an error or warning.

```
meansd <- function(x, na.rm = FALSE) {  
  if (!is.numeric(x)) {  
    message("non-numeric input detected")  
    return(NA)  
  }  
  # output  
  c(mean = mean(x, na.rm = na.rm),  
    sd = sd(x, na.rm = na.rm))  
}
```

# Generic Message

```
# no message
```

```
meansd(c(4, 5, 3, 1, 2))
```

```
      mean      sd  
3.000000 1.581139
```

```
# message
```

```
meansd(c(TRUE, FALSE, TRUE, FALSE))
```

```
non-numeric input detected
```

```
[1] NA
```

# Documenting Functions

# Documenting Functions

Documenting a function involves adding descriptions for what the purpose of a function is, the inputs it accepts, and the output it produces.

- ▶ Description: what the function does
- ▶ Input(s): what are the inputs or arguments
- ▶ Output: what is the output (returned value)

# Documenting Functions

There are several approaches for writing documentation of a function.

I will show you various examples for documenting a function.

In particular, I will show you how to use what are called **roxygen comments** to achieve this task.

While not used by most useRs, roxygen comments are great when you want to take your code and make a package out of it.

# Documenting Functions

One option to document a function is by simply adding a short description of what the arguments should be like. In this case, the description is outside the function

```
# function for adding two numbers  
# x: number  
# y: number  
add <- function(x, y) {  
  x + y  
}
```



# Documenting Functions

In this case, the description is between `<-` and `function()`

```
add <-  
  # function for adding two numbers  
  # x: number  
  # y: number  
  function(x, y) {  
    x + y  
  }
```

# Documenting Functions

In this case, the description is inside the function

```
add <- function(x, y) {  
  # function for adding two numbers  
  # x: number  
  # y: number  
  x + y  
}
```

# Documenting Functions

In this case, the description is inside the function

```
# description of arguments
compound_interest <- function(principal = 1, rate = 0.01,
                               periods = 1, time = 1)
{
  # principal = Principal Amount
  # rate = Annual Nominal Interest Rate as a decimal
  # time = Time Involved in years
  # periods = number of compounding periods per unit time
  principal * (1 + rate/periods)^(time * periods)
}
```

# Roxygen Comments

One interesting option to document functions is by using **roxygen comments**

```
#' @title Standardize  
#' @description Transforms values in standard units  
#' @param x numeric vector  
#' @param na.rm whether to remove missing values  
#' @return standardized values  
#' @examples  
#'   standardize(runif(10))  
standardize <- function(x, na.rm = FALSE) {  
  z <- (x - mean(x, na.rm = na.rm)) / sd(x, na.rm = na.rm)  
  return(z)  
}
```

# Roxygen Comments

Roxygen comments are R comments formed by the hash symbol immediately followed by an apostrophe: `#'`

You specify the label of a field with `@` and a keyword: e.g. `@title`

The syntax highlighting of RStudio recognizes this type of comments and labels

# Typical roxygen fields

label	meaning	description
@title	title	name of your function
@description	description	what the function does
@param	parameter	describe input parameter
@return	output	what is the returned value
@example	example	one or more usage examples

Time permitting, at the end of the semester we'll see how using Roxygen comments simplifies the creation of an R package.

# General Recommendations

# On Writing Functions

Before you write a function:

- ▶ Write down the goal. What should the function do?
  - Draw a picture if it helps clarify the goal
- ▶ Check whether the function already exists.
  - Check base R, packages, code you've written, or google it
- ▶ Write down the inputs and outputs
- ▶ Write code to handle a simple case



# On Writing Functions

- ▶ Baby steps: always start simple with toy-values or small data sets
- ▶ Work first on what will be the body of the function
- ▶ Check out each step of the way (don't worry yet about efficiency or elegance or cleverness)
- ▶ Don't try to do much at once
- ▶ Create the function (i.e. encapsulate the body) once everything works
- ▶ Don't write long functions: write short / small functions (preferably less than 10 lines of code)
- ▶ Personally, I think more than 20 lines of code are too many.

# On Writing Functions

- ▶ Include documentation; we suggest using Roxygen comments.
- ▶ Optional: after you have a function that works, then you may worry about “elegance”, “efficiency”, “cleverness”, etc.
- ▶ Often, it’s better to have an “ugly/inefficient” function that does the work, rather than wasting a lot of time, effort, and energy to get a “smart” function.
- ▶ The more you practice, the easier will be to create functions.
- ▶ As you get more experience, making more clever and elegant functions will be less difficult, and worth your time.