

# Performance and Profiling

Stat 33B: Intro to Advanced Programming in R

Gaston Sanchez

Notes 13

# About

In this lecture we will talk about performance in R, describing a simple example from different approaches, measuring their performance, and profiling them to identify bottlenecks in the code. Overall, we touch on:

- ▶ Copy-on-modify policy
- ▶ Performance
- ▶ Profiling

```
# suggested packages  
library(lobstr)  
library(microbenchmark)  
library(profvis)
```

# Motivation

# Toy Example

Consider this example:

$$\begin{array}{c} \mathbf{x} \\ \boxed{2} \boxed{4} \boxed{6} \boxed{8} \end{array} + \begin{array}{c} \mathbf{y} \\ \boxed{0} \boxed{1} \boxed{0} \boxed{1} \end{array} = \begin{array}{c} \mathbf{z} \\ \boxed{2} \boxed{5} \boxed{6} \boxed{9} \end{array}$$

Let's approach this with various coding options

# Toy Example

<b>x</b>					<b>y</b>					<b>z</b>			
2	4	6	8	+	0	1	0	1	=	2	5	6	9

## A) For Loop 1

```
x = c(2,4,6,8)
y = c(0,1,0,1)
z = c()

for (i in 1:length(x)) {
  z = c(z, x[i] + y[i])
}
```

## B) For Loop 2

```
x = c(2,4,6,8)
y = c(0,1,0,1)
z = rep(0, length(x))

for (i in 1:length(x)) {
  z[i] = x[i] + y[i]
}
```

## C) Vectorized

```
x = c(2,4,6,8)
y = c(0,1,0,1)
z = x + y
```

What does R do in each case?

# Measuring Performance with `system.time()`

## Measuring Performance: case A)

```
# a) For Loop 1
x <- runif(10000)
y <- runif(10000)
z <- c()

case_a <- system.time({
  for (i in seq_along(x)) {
    z <- c(z, x[i] + y[i])
  }
})
case_a
```

```
##      user  system elapsed
##    0.180    0.128    0.314
```

## Measuring Performance: case B)

```
# b) For Loop 2
x <- runif(10000)
y <- runif(10000)
z <- rep(0, 10000)

case_b <- system.time({
  for (i in seq_along(x)) {
    z[i] <- x[i] + y[i]
  }
})
case_b
```

```
##      user  system elapsed
##    0.004    0.000    0.004
```



## Measuring Performance: case C)

```
# c) Vectorized code
x <- runif(10000)
y <- runif(10000)

case_c <- system.time({
  z <- x + y
})
case_c
```

```
##      user  system elapsed
##         0         0         0
```

## Comparison

```
rbind(case_a, case_b, case_c)
```

	user.self	sys.self	elapsed	user.child	sys.child
case_a	0.180	0.128	0.314	0	0
case_b	0.004	0.000	0.004	0	0
case_c	0.000	0.000	0.000	0	0

- ▶ As you can tell, vectorized code (C) beats both for() loops
- ▶ Moreover, case (A) is much slower than case (B)
- ▶ What's going on in each case?

# Performance comparison

- ▶ Conceptually, all three options (A, B, C) are equivalent.
- ▶ In R, however, they all have different performances.
- ▶ To see why this is, we'll use functions from packages:
  - "lobstr"
  - "microbenchmark"

# Motivation Example

<b>x</b>					<b>y</b>					<b>z</b>			
2	4	6	8	+	0	1	0	1	=	2	5	6	9

## A) For Loop 1

```
x = c(2,4,6,8)
y = c(0,1,0,1)
z = c()

for (i in 1:length(x)) {
  z = c(z, x[i] + y[i])
}
```

## B) For Loop 2

```
x = c(2,4,6,8)
y = c(0,1,0,1)
z = rep(0, length(x))

for (i in 1:length(x)) {
  z[i] = x[i] + y[i]
}
```

## C) Vectorized

```
x = c(2,4,6,8)
y = c(0,1,0,1)
z = x + y
```

## Discussion example (A)

```
x = c(2, 4, 6, 8)
y = c(0, 1, 0, 1)
z = c()

for (i in 1:length(x)) {
  z = c(z, x[i] + y[i])
}
```

- ▶ z has to grow at every iteration
- ▶ It seems that there are only two functions in the loop: `+` and `c()`
- ▶ In fact, `for()` is a function, as well as `:"`, `length()`, `z[i]=`, `x[i]`, `y[i]`, `+"`

## Discussion example (B)

```
x = c(2, 4, 6, 8)
y = c(0, 1, 0, 1)
z = rep(0, 4)

for (i in 1:length(x)) {
  z[i] = x[i] + y[i]
}
```

- ▶ here, z does not have to grow at every iteration
- ▶ There are still several function calls: `for()`, `":"`, `length()`, `x[i]`, `y[i]`, `+`, `=`

## Discussion example (C)

```
x = c(2, 4, 6, 8)
y = c(0, 1, 0, 1)
z = x + y
```

- ▶ In this case, to obtain `z` there are only two function calls: `"+"` and `"="`
- ▶ Both are **primitive** functions which means they are written in C

# Profiling



# Profiling

- ▶ To understand code performance we use a **profiler**
- ▶ R uses a fairly simple type called a *sampling* or *statistical profiler*
- ▶ A sampling profiler stops the execution of code every few milliseconds and records the call stack
- ▶ The default profiling resolution is quite small, so if your function takes even a few seconds it will generate hundreds of samples
- ▶ We'll use the "profvis" package to visualise aggregates

## Profiling with Rprof(): example (A)

```
x = runif(100000)
y = runif(100000)
z = c()

Rprof()
invisible(
  for (i in 1:length(x)) {
    z = c(z, x[i] + y[i])
  }
)
Rprof(NULL)
```

Note: I've increased the size of x and y for profiling purposes.

## Profiling with Rprof(): example (A)

```
summaryRprof()
```

```
$by.self
```

	self.time	self.pct	total.time	total.pct
"c"	25.76	100	25.76	100

```
$by.total
```

	total.time	total.pct	self.time	self.pct
"c"	25.76	100	25.76	100

```
$sample.interval
```

```
[1] 0.02
```

```
$sampling.time
```

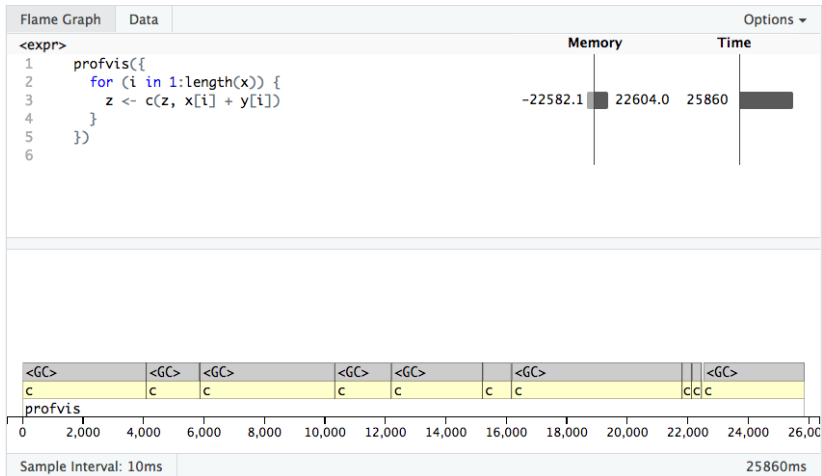
```
[1] 25.76
```

## Profiling with profvis(): Example A

```
x = runif(100000)
y = runif(100000)
z = c()

profvis({
  for (i in 1:length(x)) {
    z <- c(z, x[i] + y[i])
  }
})
```

# Profiling with profvis(): Example A (cont'd)



## Profiling with `profvis()`: Example A (cont'd)

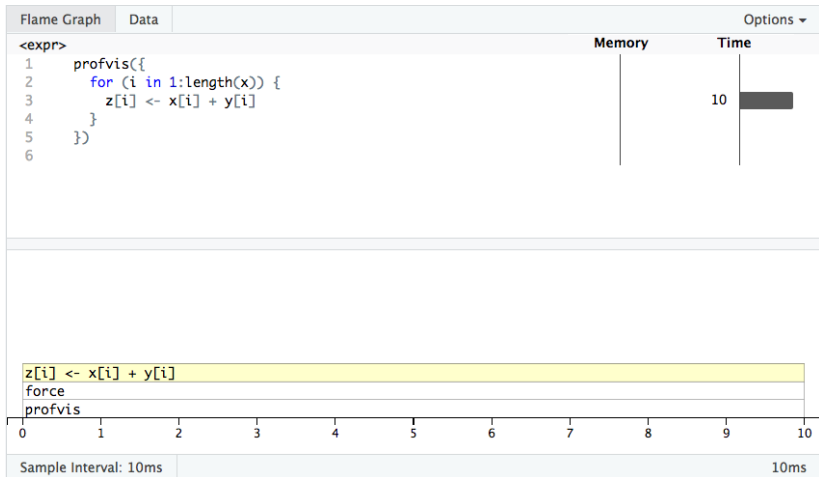
- ▶ The top pane shows the source code, overlaid with bar graphs for memory and execution time for each line of code.
- ▶ This display gives you a good overall feel for the bottlenecks. You can see that the whole process takes about 25.86 seconds.
- ▶ The bottom pane displays a flame graph showing the full call stack. This allows you to see that `c()` is constantly called.
- ▶ GC indicates that the garbage collector is running. Here, GC is taking a lot of time, which is usually an indication that we're creating many short-lived objects.

## Profiling with profvis(): Example B

```
x = runif(100000)
y = runif(100000)
z = rep(0, 100000)

profvis({
  for (i in 1:length(x)) {
    z[i] <- x[i] + y[i]
  }
})
```

## Profiling with profvis(): Example B (cont'd)



As you can tell, example (B) “only” takes about 10 milliseconds



## Profiling with `profvis()`: Example C

```
x = runif(100000)
y = runif(100000)

profvis({
  z <- x + y
})
```

This will fail because the code runs too fast for the profiler. And also because the profiling tool in R does not extend to C code.

# Microbenchmarks

# Microbenchmarking

- ▶ A microbenchmark is a measurement of the performance of a very small piece of code.
- ▶ We are going to use microbenchmarks (from the package "microbenchmark") to illustrate the performance of various pieces of R code.
- ▶ Keep in mind that we will review simple examples so that you gain intuition of how R works.
- ▶ In general, don't change the way you code because of these microbenchmarks.

# Microbenchmarking

Here's an example comparing two ways of calculating square root:

`sqrt(x)` against `x^0.5`

```
library(microbenchmark)
```

```
x <- runif(100)
```

```
microbenchmark(
```

```
  sqrt(x),
```

```
  x^0.5
```

```
)
```

```
## Unit: nanoseconds
```

```
##      expr   min      lq    mean median    uq   max neval
```

```
##  sqrt(x)   315   336.0  634.71    370   425 19876   100
```

```
##   x^0.5 2233 2260.5 2940.36   2314 2382 18153   100
```

# Microbenchmarking

- ▶ By default, `microbenchmark()` runs each expression 100 times
- ▶ It randomizes the order of the expressions
- ▶ It summarizes the results with a minimum (`min`), lower quartile (`lq`), median, upper quartile (`uq`), and maximum (`max`)
- ▶ Focus on the median, and the quartiles to get a feel for the variability

# About Performance in R

# Performance of R

- ▶ R is not a particularly fast language.
- ▶ This is not an accident but mainly a consequence of R being inspired by S.
- ▶ R (like S) has been designed to make data analysis and statistics easier for you to do, not to make life easier for your computer.
- ▶ By design, R privileges flexibility and convenience (for you) over performance (for your computer).
- ▶ For most applications and purposes, R is fast enough.
- ▶ After all, for really high-performance code, you must always use languages that are closer to the computer (e.g. C, C++, Fortran).

# Why is R slow?

Some of the main reasons why R can be slow:

- ▶ Because of poorly written code.
- ▶ Because R is an extremely dynamic programming language.
- ▶ Because of its lexical scoping rules—together with extreme dynamism.
- ▶ Because of lazy evaluation overhead.



Some of the reasons why R  
is not particularly fast

# No scalars in R

- ▶ There are no scalars in R
- ▶ This is intentional: because S didn't have scalars too

This has implications because something apparently simple as

2 + 3

involves adding two vectors instead of adding two scalars.

- ▶ For most applications, though, I don't think this is a serious concern.

# Extreme Dynamism

- ▶ R is an extremely dynamic programming language
- ▶ This means that you can change almost any object.
- ▶ This is the reason why you can write code like this:

```
a = TRUE    # start with a logical vector

a[2] = 3L   # add a second element, and change 'a' into integer

a[1] = -2   # modify 1st elem, and change 'a' into double

a = as.character(a) # coerce it into character

# make it into a function, or 'closure'
a = function(a) {
  a*a
}
```

# Extreme Dynamism

## Pros:

- ▶ you need minimal upfront planning
- ▶ gives you plenty of flexibility
- ▶ you can change your mind at any time, iterating your way to a solution without having to start afresh

## Cons:

- ▶ too much flexibility has a cost for the language interpreter
- ▶ it's difficult to predict exactly what will happen with a given function call
- ▶ R has to constantly look for the right names and their methods

## Name lookup with mutable environments

In this example, each time we print a it comes from a different environment:

```
a = 1

f = function() {
  g = function() {
    print(a)
    assign("a", 2, envir = parent.frame())
    print(a)
    a = 3
    print(a)
  }
  g()
}
f()
```

## Name lookup with mutable environments

- ▶ It's surprisingly difficult to find the value associated with a name in R
- ▶ This is due to combination of lexical scoping and extreme dynamism
- ▶ This means that R can't do name lookup just once, it has to start from scratch each time
- ▶ This issue is accentuated by the fact that almost every operation is a lexically scoped function call

## Name lookup with mutable environments

```
make_taco = function(tortilla, carne) {  
  (tortilla + carne)^2  
}
```

- ▶ It seems that `make_taco()` calls two functions: `+` and `^`
- ▶ In fact, it calls four because `{` and `(` are regular functions in R

# Lazy evaluation overhead

- ▶ Lazy evaluation means that function arguments are evaluated when needed. In other words, if an argument is not used, then it won't be evaluated.
- ▶ To implement lazy evaluation, R uses a **promise** object that contains the expression needed to compute the result and the environment in which to perform the computation.
- ▶ Creating promise objects has some overhead, so each additional argument to a function decreases its speed a little.



# Lazy evaluation overhead

Consider this example in which each version of the function has one additional argument:

```
g0 = function() NULL
```

```
g1 = function(a = 1) NULL
```

```
g2 = function(a = 1, b = 2) NULL
```

```
g3 = function(a = 1, b = 2, c = 3) NULL
```

```
g4 = function(a = 1, b = 2, c = 3, d = 4) NULL
```

```
g5 = function(a = 1, b = 2, c = 3, d = 4, e = 5) NULL
```

# Lazy evaluation overhead

```
microbenchmark(g0(), g1(), g2(), g3(), g4(), g5(), times = 10000)
```

```
## Unit: nanoseconds
```

##	expr	min	lq	mean	median	uq	max	neval
##	g0()	108	118	254.5652	124	133	929729	10000
##	g1()	155	170	282.2761	178	191	444477	10000
##	g2()	195	215	349.4013	226	244	452049	10000
##	g3()	237	261	425.8221	274	302	407385	10000
##	g4()	276	300	473.1978	315	360	408547	10000
##	g5()	313	344	573.9759	362	425	433378	10000

# Various Examples

## Example

```
# median-centering a data frame
dat = data.frame(matrix(runif(100 * 1000), ncol = 100))
medians = vapply(dat, median, numeric(1))

for (j in seq_along(medians)) {
  dat[,j] = dat[,j] - medians[j]
}
```

- ▶ every iteration of the loop copies the data frame
- ▶ this is because `[<-data.frame` is not a primitive function
- ▶ so every replacement produces an entire copy

## Example

```
# median-centering a data frame
```

```
dat = data.frame(matrix(runif(100 * 1000), ncol = 100))
```

```
medians = vapply(dat, median, numeric(1))
```

```
lis = as.list(dat)
```

```
for (j in seq_along(medians)) {  
  lis[[j]] = lis[[j]] - medians[j]  
}
```

- ▶ using a list instead of data frame improves performance
- ▶ having a list, we can use `[[<-list` which is a primitive function
- ▶ here, every replacement occurs in place (without any copies)

# Some Comments

# My Advice

- ▶ **Good code is clear:** it's better to be clear than clever
- ▶ **Time and practice:** it takes time and practice to be able to write code in R that performs well, so don't obsess too much about it (if you don't need to)
- ▶ **Is it worth your time?** Don't spend hours of your time to save seconds of computer time when you can get the job done without worrying about performance.
- ▶ **Know when to quit:** At some point, if you are seriously concerned about the speed of your code, you will reach a limit within R, and will have to switch to “faster” languages (e.g. C, C++, Fortran)

# My Advice

- ▶ Start small (baby steps), writing code that gets the job done
- ▶ Gradually work your way up, developing and testing your code under different circumstances
- ▶ The more you understand R, you'll tend to have more opportunity to improve your code
- ▶ The more time you have to think about your code (special cases, make it more robust, more user friendly), the higher the chances to make it faster
- ▶ What's the ROI? Sometimes you don't have time to polish your code, so even if you know your code is not optimized, this won't matter because you need to move to next stage(s) of your project(s)